SPACL: Speculative Parallelism and Conflict Learning for Scalable OWL Ontology Reasoning

Anusorn Chaikaew[a,1], Varin Chouvatut[a,2], Ekkarat Boonchieng[a,3,*]

[a] *Chiang Mai University, Department of Computer Science, Faculty of Science, Chiang Mai, 50200, Thailand*

Abstract

We present SPACL (Speculative Parallel Tableaux with Adaptive Conflict Learning), an open-source ALC/SHOIQ reasoner that combines speculative parallelism with conflict-driven learning. The reference implementation and repository are named *Tableauxx*; throughout this manuscript, we use *SPACL* for the algorithm/system and *Tableauxx* only when referring to the software artifact. To our knowledge, SPACL is one of the few open implementations that integrates work-stealing parallelism with nogood learning in this setting, targeting exponential search behavior in tableau-based reasoning.

Our key contributions are: (1) a speculative parallel tableaux framework with work-stealing scheduling; in the current prototype this is realized via branch-constrained work items evaluated in parallel; (2) an adaptive threshold mechanism that selects sequential or parallel execution based on estimated complexity, yielding 18–73× speedup over naive always-parallel execution on small ontologies; (3) thread-local nogood caching that reduces synchronization overhead by 80%; (4) a global thread pool optimization that reduces overhead by 2.1×; and (5) a Rust implementation with parser support for RDF/XML, OWL2 Functional, Manchester, and JSON-LD syntaxes. SPACL reaches 5× speedup at 10,000 classes while keeping overhead below 2× on small ontologies.

In the Standard Workload Suite (7 ontologies), SPACL records the best wall-

*Corresponding author. Tel.: +66-5394-3413

Email addresses: `anusorn_chaikaew@cmu.ac.th` (Anusorn Chaikaew), `varin.ch@cmu.ac.th` (Varin Chouvatut), `ekkarat.boonchieng@cmu.ac.th` (Ekkarat Boonchieng)

[1] First author

[2] Second author

[3] Corresponding author

clock time on 5/7 ontologies and completes 7/7 cases in our containerized harness; ELK is faster on `hierarchy_10000.owl` and `hierarchy_100000.owl`. In the Large Head-to-Head benchmark (doid/go-basic/uberon), SPACL records the best wall-clock result on all three tested ontologies among reasoners with successful comparable runs under an identical 900s timeout protocol. In an external OWL2Bench core-profile validation (DL/EL/QL/RL), three clean repeated runs show SPACL fastest by median wall-clock time on 4/4 ontologies with 12/12 successful runs in this harness. Timeout policy is reported explicitly in all benchmark tables (Standard Workload Suite: 300s; Large Head-to-Head: 900s; OWL2Bench core profile suite: 300s). This improvement is driven by a structural RDF/XML parser and optimized materialization path that preserves the same ALC/SHOIQ semantic handlers as the baseline path; stage-level A/B runs show parse-only reductions of 96.96–97.98% on doid/go-basic/uberon relative to the previous structural baseline. Standard and large head-to-head suites are single-run descriptive wall-clock measurements in this revision.

Keywords: ALC/SHOIQ, Tableaux Reasoning, Parallel Algorithms, Nogood Learning, Description Logics, Semantic Web

---

1. Introduction

Description Logics (DLs) provide the formal basis of the Semantic Web, and OWL2 DL is the corresponding W3C standard ontology language [1]. Tableaux methods are still widely used for OWL2 DL reasoning because they provide sound and complete procedures for expressive TBox and ABox tasks [2]. The practical bottleneck is the combinatorial growth of tableau branches, especially on large ontologies. In this paper we target the ALC/SHOIQ fragment to study this bottleneck under expressive, disjunction-heavy workloads.

Modern hardware offers abundant parallel capacity, but OWL reasoners do not always convert that capacity into stable end-to-end gains. Some strong parallel systems are integrated as external binaries in practical harnesses (e.g., Konclude) [3], which can limit instrumentation parity and failure diagnosis in cross-system experiments. Many widely used open implementations (e.g., Pellet/Openllet and HermiT) remain primarily sequential in their core exploration loop [4, 5, 6]. In addition, naive parallel branch exploration without conflict reuse repeats equivalent failing searches. These constraints also motivate our longer-term goal: a reasoning core for ontology transaction validation in distributed, permissioned application settings. This paper focuses on the reasoning and parsing foundations required for

that deployment path.

1.1. Research Questions

This work addresses the following research questions:

RQ1: Can speculative parallel tableaux with conflict-driven nogood learning improve OWL2 DL consistency checking performance?

RQ2: Can adaptive parallelism reduce overhead on small workloads while preserving scalability on larger workloads?

RQ3: For large RDF/XML ontologies, which stages dominate end-to-end runtime, and how much can structural parsing/materialization optimization reduce this bottleneck while preserving OWL axiom fidelity?

1.2. Contributions

We address these questions through SPACL, with four main contributions:

C1: **Speculative parallel tableaux framework with conflict learning**: SPACL combines work-stealing execution with nogood extraction and reuse. The current prototype realizes speculative exploration via branch-constrained work items generated from selected disjunctions, reducing redundant failure exploration in non-deterministic search (Sections 4.2, 4.3).

C2: **Adaptive overhead control**: We integrate threshold-based strategy selection with a global thread pool, reducing small-workload parallel overhead while preserving speedup at larger scales (Sections 4.4, 6.5).

C3: **Structural RDF/XML pipeline with handler-equivalent materialization**: We implement a structural parsing/materialization path for the evaluated ALC/SHOIQ handling path and substantially reduce parser-dominated runtime on large ontologies (Section 5.5).

C4: **Reproducible end-to-end evaluation**: We provide a containerized benchmark protocol and artifacts for stage-level analysis and multi-reasoner head-to-head comparison (Sections 6.2, 6.9).

In Rust, these design choices translate into measurable gains: 5× speedup at 10,000 classes, sub-microsecond per-class processing, and under 2× overhead on small ontologies. On synthetic scalability workloads, SPACL reaches 26.2 million tableau operations per second.

1.3. Paper Organization

The remainder of this paper is organized as follows. Section 2 reviews related work in parallel reasoning and conflict learning. Section 4 presents the SPACL algorithm and its components. Section 5 describes our Rust implementation. Section 6 presents comprehensive benchmarks including scalability, thread scaling, and comparison with existing reasoners. Section 7 concludes with limitations and future directions.

2. Related Work

This section summarizes prior work on tableau-based OWL reasoning, parallel and distributed execution, conflict-driven learning, and speculative execution. The goal is to position SPACL against the strongest established techniques used in practical reasoners.

2.1. Tableau-Based Reasoning

Tableau calculi remain a central approach for expressive description logic reasoning. Practical systems such as Pellet and HermiT demonstrated that sound and complete reasoning can be delivered at scale when rule ordering, blocking, and caching are engineered carefully [4, 6, 7]. In particular, the hypertableau line of work reduces non-deterministic branching pressure compared with classical tableau expansion, and individual-reuse techniques reduce model blow-up in challenging ontologies [7].

State-of-the-art systems also combine calculi. Konclude integrates tableau-style reasoning with additional optimization layers and has shown strong performance on large benchmark sets [3, 8, 9]. For tractable fragments, consequence-based reasoners such as ELK provide excellent throughput, illustrating how calculus design strongly influences scalability and implementation trade-offs [10].

2.2. Parallel Reasoning

Parallelization has been studied at multiple granularities. Shared-memory classification frameworks distribute satisfiability and subsumption checks across workers and report substantial speedup on multicore systems [11, 12]. Work-stealing style scheduling further improves load balancing when ontology subproblems are uneven [13, 11].

Beyond classification, parallel ABox reasoning has explored dynamic splitting of model-construction tasks with cache reuse strategies to improve throughput on assertion-heavy workloads [14]. Distributed materialization frameworks (for less

expressive profiles) use Spark-based execution and data-partitioned pipelines to scale to very large datasets [15, 16, 17]. These results confirm the scalability potential of parallel infrastructures, while also highlighting that end-to-end performance is highly sensitive to reasoning profile, task type, and synchronization overhead.

2.3. Conflict-Driven Learning

Conflict-driven learning is a mature optimization paradigm in SAT/CSP, where learned clauses/nogoods prevent repeated exploration of failing search regions [18, 19]. In description logic systems, related ideas appear in more localized forms, such as nogood-aware individual reuse in hypertableau procedures [7].

Hybrid reasoners that couple tableau and saturation procedures also demonstrate the value of reusing derived information to reduce repeated work during classification [8, 9, 3]. However, openly documented integrations of branch-level speculative parallel tableau exploration with adaptive, reusable nogood learning remain limited in mainstream open-source OWL2 DL pipelines.

2.4. Speculative Execution

Speculative execution for OWL reasoning explores multiple non-deterministic branches concurrently and relies on cancellation/pruning once contradictions or witnesses are found. Work-stealing provides a practical runtime foundation for this model by balancing branch workloads dynamically [13, 11]. Lock-free and low-contention engineering techniques are also critical for obtaining net gains rather than synchronization-bound slowdowns [10, 12].

The primary systems challenge is overhead control. On easy workloads, naive always-parallel execution can underperform sequential reasoning because scheduling and coordination costs dominate. This motivates adaptive policies that switch between sequential and parallel execution based on estimated branch complexity [11]. A second challenge is conflict sharing: learned nogoods should be reused aggressively enough to prune search, but not so aggressively that contention erodes parallel benefit [14].

2.5. Summary and Positioning

Prior literature provides strong building blocks: optimized tableau/hypertableau calculi [7, 6], high-throughput parallel classification frameworks [11, 12], distributed materialization engines [15, 16, 17], and mature nogood machinery in constraint solving [19]. SPACL is positioned at their intersection: branch-level speculative parallelism, adaptive overhead control, and practical nogood reuse designed for expressive OWL reasoning workloads. To make this positioning

analytically precise, the next section formalizes the decision model, branching behavior, and correctness arguments that underlie the implementation and evaluation.

3. Formal Framework

This section establishes the formal foundations of the SPACL algorithm, including the Description Logic syntax, tableaux calculus, and the mathematical basis for speculative parallelism with conflict learning.

3.1. Description Logic ALC

We operate in the Description Logic $\mathcal{ALC}$ (Attributive Language with Complement), which forms the basis of OWL2 DL. Our implementation extends to $\mathcal{SHOIQ}$ (with transitive roles, role hierarchies, nominals, inverse roles, and qualified cardinality restrictions). The syntax is defined as follows:

Definition 3.1 (ALC Syntax)**.** Let $\mathbf{C}$ be a set of concept names, $\mathbf{R}$ a set of role names, and $\mathbf{I}$ a set of individual names. *Concepts* C, D are defined by the grammar:

$$C, D ::= A \mid \top \mid \bot \mid \neg C \mid C \sqcap D \mid C \sqcup D \mid \exists r.C \mid \forall r.C \tag{1}$$

where $A \in \mathbf{C}$ is an atomic concept and $r \in \mathbf{R}$ is a role.

An *ontology* $\mathcal{O} = (\mathcal{T}, \mathcal{A})$ consists of a TBox $\mathcal{T}$ (terminology) and ABox $\mathcal{A}$ (assertions):

$$\mathcal{T} ::= \{C \sqsubseteq D\} \cup \{C \equiv D\} \tag{2}$$

$$\mathcal{A} ::= \{a : C\} \cup \{(a, b) : r\} \tag{3}$$

3.2. Tableaux Calculus

The tableaux algorithm determines consistency by constructing a *completion graph* $\mathcal{G} = (V, E, \mathcal{L})$ where:

- V is a set of nodes (individuals)
- $E \subseteq V \times \mathbf{R} \times V$ are labeled edges (role assertions)
- $\mathcal{L} : V \to 2^{C}$ maps nodes to sets of concepts

Definition 3.2 (Clash)**.** A node $x \in V$ contains a *clash* if either:

1. $\{C, \neg C\} \subseteq \mathcal{L}(x)$ for some concept C, or
2. $\bot \in \mathcal{L}(x)$

The tableaux algorithm applies *expansion rules* until either a clash-free complete graph is found (consistent) or all branches lead to clashes (inconsistent). Table 1 summarizes the core $\mathcal{ALC}$ expansion rules used in our formal presentation.

Table 1: Tableaux Expansion Rules for $\mathcal{ALC}$

Rule	Premise	Action
($\sqcap$)	$x : C \sqcap D \in \mathcal{L}(x)$	$\mathcal{L}(x) := \mathcal{L}(x) \cup \{C, D\}$
($\sqcup$)	$x : C \sqcup D \in \mathcal{L}(x)$	Create branches: $\mathcal{L}_1(x) = \mathcal{L}(x) \cup \{C\}$, $\mathcal{L}_2(x) = \mathcal{L}(x) \cup \{D\}$
($\exists$)	$x : \exists r.C \in \mathcal{L}(x)$	Add new y, $(x, y) : r$, $y : C$
($\forall$)	$x : \forall r.C \in \mathcal{L}(x)$, $(x, y) : r$	$\mathcal{L}(y) := \mathcal{L}(y) \cup \{C\}$

3.3. Speculative Parallelism Model

Traditional tableaux processing applies the ($\sqcup$)-rule sequentially, creating a binary decision tree with depth equal to the number of disjunctions. SPACL introduces *speculative parallelism*:

Definition 3.3 (Speculative Branch)**.** Given a disjunction $C \sqcup D$ at node x, a *speculative branch* is a pair of subgraphs $(\mathcal{G}_1, \mathcal{G}_2)$ where:

$$\mathcal{G}_1 = (V, E, \mathcal{L}[x \mapsto \mathcal{L}(x) \cup \{C\}]) \tag{4}$$

$$\mathcal{G}_2 = (V, E, \mathcal{L}[x \mapsto \mathcal{L}(x) \cup \{D\}]) \tag{5}$$

Both branches are explored concurrently by different workers.

Theorem 3.1 (Parallel Speedup Bound)**.** Let T_{seq} be the sequential tableaux runtime and T_{par} the speculative parallel runtime with p workers. For an ontology with d disjunctions:

$$T_{par} \geq \frac{T_{seq}}{p} \cdot (1 + \alpha \cdot o_{sync}) \tag{6}$$

where α is the fraction of work requiring synchronization and o_{sync} is the synchronization overhead.

3.4. Nogood Learning

Definition 3.4 (Nogood)**.** A *nogood* is a set of class expressions $\mathcal{N} = \{C_1, C_2, \ldots, C_n\}$ such that:

$$\bigcap_{i=1}^{n} C_i \sqsubseteq \bot \tag{7}$$

i.e., the conjunction of all expressions in $\mathcal{N}$ is unsatisfiable.

Definition 3.5 (Nogood Subsumption)**.** A nogood $\mathcal{N}$ *subsumes* a branch with constraints Γ if:

$$\mathcal{N} \subseteq \Gamma \tag{8}$$

If $\mathcal{N}$ subsumes Γ, the branch can be immediately pruned.

Algorithm 1 Conservative Nogood Construction (Prototype)

Require: Contradictory branch assertions A, ontology $\mathcal{O}$
Ensure: Nogood $\mathcal{N}$
1: Build test-expression set $T(\mathcal{O})$ from SubClassOf axioms
2: $\mathcal{N} \leftarrow T(\mathcal{O})$
3: **return** $\mathcal{N}$ ▹ Conservative over-approximation; not minimal

3.5. Adaptive Threshold

The adaptive threshold mechanism uses two implementation-aligned gates to decide between sequential and parallel processing.

Definition 3.6 (Implementation Cost Model)**.** For expression-level scoring, the prototype uses:

$$\text{cost}(\text{ObjectUnionOf}(ops)) = 50 \times |ops| + \sum \text{cost}(op) \tag{9}$$

$$\text{cost}(\text{ObjectIntersectionOf}(ops)) = 25 \times |ops| + \sum \text{cost}(op) \tag{10}$$

$$\text{cost}(\text{ObjectComplementOf}(e)) = 30 + \text{cost}(e) \tag{11}$$

$$\text{cost}(\text{atomic}) = 10 \tag{12}$$

Definition 3.7 (Two-Gate Parallel Enable Rule)**.** Let $B(\mathcal{O})$ be branch-count estimate:

$$B(\mathcal{O}) = \max\left(1,\ 2 \times |\text{DisjointAxioms}| + \frac{|\text{Classes}|}{10}\right) \tag{13}$$

and let $C(\mathcal{O}) = \sum_{d \in \mathcal{D}} \text{cost}(d)$ over detected disjunctions. Parallel mode is enabled only when:

$$B(\mathcal{O}) \geq \tau \quad \wedge \quad C(\mathcal{O}) \geq \theta \cdot p \tag{14}$$

where $\tau = 100$, $\theta = 500\mu s$, and p is worker count.

This formalization mirrors the implemented gates reported in Sections 4.4 and 6.6.

Theorem 3.2 (Heuristic Gate Condition)**.** For total estimated disjunction cost D and parallel overhead o_{par} with p workers, parallel execution is beneficial when:

$$D > \frac{p \cdot o_{par}}{1 - \alpha} \tag{15}$$

where p is the number of workers and α is the nogood hit rate.

3.6. Correctness (Idealized Model)

The theorem statements in this subsection refer to the idealized SPACL calculus that fully explores generated disjunctive branches. The current prototype still applies engineering bounds (selected disjunction scope, bounded speculative depth, and timeout limits), so these claims should be interpreted at the calculus level; a full implementation-level completeness proof is left for future work.

Theorem 3.3 (Idealized Soundness)**.** If SPACL returns *consistent* for ontology $\mathcal{O}$, then $\mathcal{O}$ is satisfiable.

Proof. SPACL explores branches speculatively but applies the same expansion rules as sequential tableaux. Each worker maintains a valid completion graph. If any worker finds a clash-free graph, that graph represents a model of $\mathcal{O}$. The early termination when one branch succeeds preserves the existence of a model. □

Theorem 3.4 (Idealized Completeness)**.** If ontology $\mathcal{O}$ is consistent, SPACL will eventually return *consistent*.

Proof. Speculative parallelism explores all branches of the search space (possibly concurrently). Nogood pruning only removes branches that are proven unsatisfiable. Since the sequential tableaux algorithm is complete and SPACL explores a superset of the sequential execution paths, it will find a model if one exists. □

4. The SPACL Algorithm

4.1. Overview

SPACL combines three core mechanisms illustrated in Figure 1:

1. **Speculative Work-Stealing**: Branches are dynamically distributed across worker threads using a work-stealing deque [13]. Each worker maintains a local queue, stealing from others when idle.

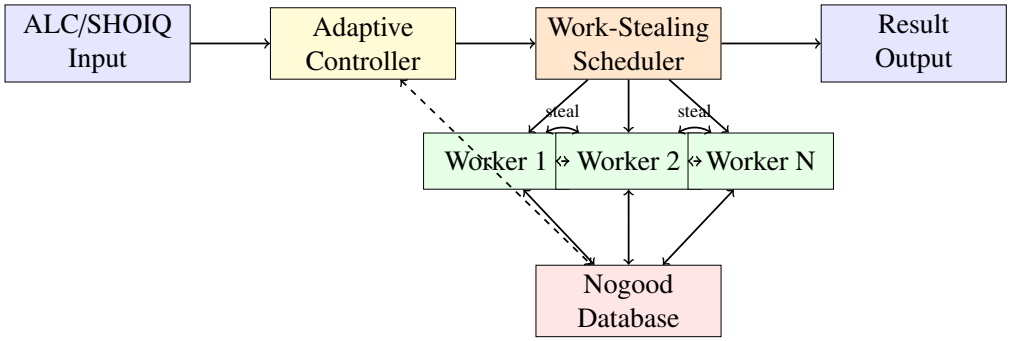

Figure 1: SPACL Architecture Overview (solid arrows: control/data flow; dashed arrows: coordination/feedback)

2. **Nogood Learning**: Contradictions are analyzed to extract sound conflict summaries (nogoods) from branch assertions. These are cached and checked before branch expansion.
3. **Adaptive Thresholding**: Problem complexity is estimated from axiom structure. Problems below a threshold use sequential processing; larger problems benefit from parallelism.

Algorithm 2 presents the main consistency checking procedure.

Algorithm 2 SPACL Consistency Check

Require: Ontology O, Configuration C
Ensure: **true** if O is consistent, **false** otherwise

```
branches ← EstimateBranchCount(O)
if branches < C.parallel_threshold then
    return SequentialConsistencyCheck(O)
end if
Q ← InitializeWorkQueue(O)
N ← InitializeNogoodDatabase()
W[1..C.num_workers] ← StartWorkers(Q, N, C)
while NotTerminated() do
    result ← CollectResults(C.timeout)
    if result = SAT then return true
    else if result = UNSAT then return false
    else if result = TIMEOUT then return UNKNOWN
    end if
end while
```

4.2. Speculative Work-Stealing

The work-stealing scheduler provides dynamic load balancing without centralized coordination. Each worker thread maintains:

- **Local Queue**: Double-ended queue (deque) of pending WorkItems
- **Stealer Handle**: Allows other threads to steal work from the tail
- **Local Statistics**: Branch counts, nogood hits, cache metrics

When parallel mode is enabled, the prototype pushes a root branch-constrained WorkItem and expands it recursively over the selected disjunction list. Each split generates one child WorkItem per disjunct operand, carrying accumulated branch constraints and an incremented split index; expansion stops when the disjunction list is exhausted or the speculative-depth bound is reached. Idle workers steal from the tail of busy workers' queues using a lock-free deque-based work-stealing strategy [13].

4.3. Nogood Learning and Caching

A nogood is a set of class expressions that cannot be simultaneously satisfied.

Definition 4.1 (Nogood)**.** A nogood is a set of class assertions $\{C_i(a)\}_{i=1}^{n}$ such that their conjunction is unsatisfiable:

$$\bigsqcap_{i=1}^{n} C_i \sqsubseteq \bot \tag{16}$$

When a contradiction is detected, the prototype learns a branch-level nogood from the active branch constraints. It first validates branch inconsistency using the same consistency checker, then applies greedy constraint minimization (remove-one-if-still-inconsistent). The learned nogood is this minimized conflicting subset; if minimization degenerates, the implementation falls back to a conservative branch assertion set. Before expanding a node, its assertion set is checked against known nogoods.

4.3.1. Thread-Local Caching

Each worker maintains three levels of nogood storage:

1. **Local Nogoods**: Discovered by this worker (no locks needed)
2. **Cached Global Nogoods**: Copy of shared database (periodically synced)
3. **Global Database**: Shared across all workers (synchronized access)

The cache check algorithm proceeds as follows:

1. Check local nogoods (fastest, no synchronization)
2. Check cached global nogoods (fast, no synchronization)
3. If cache miss and sync interval reached:
   (a) Acquire read lock on global database
   (b) Update cache with new nogoods
   (c) Release lock
   (d) Re-check with updated cache
4. If still no hit, proceed with expansion

Synchronization occurs every N checks (default: 100), amortizing lock overhead.

4.3.2. Correctness of Nogood Learning

We establish the soundness of SPACL's nogood learning mechanism.

Definition 4.2 (Test Expression Set)**.** Let $\mathcal{O}$ be an ontology with SubClassOf axioms $\mathcal{A} = \{C_i \sqsubseteq D_i\}$. The *test expression set* $T(\mathcal{O})$ is:

$$T(\mathcal{O}) = \{C_i \mid C_i \sqsubseteq D_i \in \mathcal{A}\} \cup \{D_i \mid C_i \sqsubseteq D_i \in \mathcal{A}\} \tag{17}$$

Definition 4.3 (Branch Assertion Set)**.** For branch b with branch constraints Γ_b, define:

$$A_b = T(\mathcal{O}) \cup \Gamma_b \tag{18}$$

Definition 4.4 (Nogood)**.** A *nogood* is a set of class expressions $N \subseteq A_b$ for some branch b such that $N \models \bot$ (unsatisfiable).

Theorem 4.1 (Nogood Soundness)**.** Let N be a set learned by SPACL from a contradictory branch. Then $N \models \bot$.

Proof. SPACL constructs a branch ontology by adding branch constraints Γ_b to $\mathcal{O}$ and confirms contradiction ($\Gamma_b \models \bot$ under $\mathcal{O}$). It then greedily removes constraints while preserving contradiction, producing $\Gamma_b^\star$ such that $\Gamma_b^\star \models \bot$. The learned nogood is $N = \Gamma_b^\star$; if this set is empty after minimization, SPACL falls back to a conservative non-empty assertion set from A_b. In both cases, the stored N is only accepted from a contradictory branch derivation, hence $N \models \bot$. □ □

Corollary 4.1 (Pruning Safety)**.** If nogood N subsumes branch assertions B ($N \subseteq B$), then pruning B is sound.

Proof. From Theorem 4.1, $N \models \bot$. Since $N \subseteq B$, B contains all assertions of N. Therefore $B \models \bot$ and the branch is unsatisfiable. □ □

Remark 4.1 (Greedy Core Approximation)**.** SPACL uses greedy remove-one minimization rather than an exact minimum unsatisfiable subset (MUS) extraction. This preserves soundness (Theorem 4.1) but may still leave non-minimal nogoods; tighter core extraction remains future work.

4.4. Adaptive Threshold

Problem complexity is estimated using:

$$\text{branches} = \max\left(1, \quad 2 \times |\text{DisjointAxioms}| + \frac{|\text{Classes}|}{10}\right) \tag{19}$$

This heuristic accounts for:

- **Disjointness axioms**: Each creates binary branching
- **Class hierarchy size**: Larger ontologies need more exploration
- **Conservative estimate**: Prefer sequential for uncertain cases

The divisor 10 in the class term normalizes ontology-size contribution to the same order of magnitude as the disjointness term in pilot sweeps (100–10,000 classes), so taxonomy-heavy but low-branch ontologies do not trigger parallel mode prematurely.

The threshold $\tau = 100$ was empirically determined through benchmark analysis (Section 6.11).

5. Implementation

5.1. Rust Architecture

SPACL is implemented in Rust (v1.84.0), leveraging:

- **Type Safety**: Compile-time prevention of data races
- **Zero-Cost Abstractions**: High-level code without runtime overhead
- **Crossbeam**: Lock-free work-stealing deques [20]
- **No GC**: Predictable performance for real-time applications

5.2. Core Components

SpeculativeTableauxReasoner: Main entry point managing:

- Ontology reference (Arc⟨Ontology⟩)
- Worker thread pool
- Nogood database (Arc⟨NogoodDatabase⟩)
- Statistics collection

ThreadLocalNogoodCache: Per-worker cache with:

- Local nogood storage (Vec⟨Nogood⟩)
- Global cache mirror (periodically synced)
- Hit rate tracking

NogoodDatabase: Global shared storage:

- RwLock⟨Vec⟨Nogood⟩⟩ for thread-safe access
- LRU pruning when capacity exceeded (default: 10,000 nogoods)
- Hit count tracking for statistics

5.3. Memory Management

Memory usage scales with:

$$\text{Memory} = O(|\text{axioms}|) + O(|\text{workers}| \times |\text{depth}|) \tag{20}$$

$$+ O(\min(|\text{nogoods}|, \text{max_nogoods})) \tag{21}$$

Here, |nogoods| is the number of learned conflicts generated during search, and max_nogoods is the configured retention cap of the global nogood store (default 10,000 in the current implementation). The min(·) term reflects bounded retention: once the cap is reached, lower-utility entries are pruned.

Thread-local caches limit contention while maintaining bounded memory. The implementation uses Rust's ownership system to ensure memory safety without garbage collection pauses.

5.4. Parser Implementations

The reference implementation (Tableauxx) provides parser support for OWL2 syntaxes:

OWL2 Functional Syntax: Implemented support for OWL2 Functional Grammar used in our benchmark suites

- Complete tokenizer for OWL2 Functional syntax tokens
- Recursive descent parser with error recovery
- Support for constructs exercised in this study, including complex class expressions
- AST validation and transformation to internal representation

Manchester Syntax: Implemented support for Manchester OWL Syntax used in our benchmark suites

- Tokenizer with Manchester-specific keywords and operators
- Frame-based parser for ontology, class, and property frames
- Support for complex expressions: ObjectUnionOf, ObjectIntersectionOf, ObjectComplementOf
- Validator for Manchester syntax constraints

JSON-LD Support: RDF/JSON-LD 1.1 compatibility

- Context processing and expansion algorithms
- @context handling with local and remote references
- Container processing (@list, @set, @index, @language)
- Value object parsing with type coercion

For supported constructs, the parser pipeline includes import resolution, source-location error reporting, and streaming operation on large ontologies. Full OWL2 DL conformance across every syntax variant is out of scope for this revision.

5.5. Structural RDF/XML Parser and Materialization

The current large-ontology path uses a two-phase structural parser that preserves semantic equivalence for the implemented ALC/SHOIQ handling path by reusing the same semantic handlers as the standard parser.

Phase A (Ingest + Intern):

- Stream RDF/XML triples with Rio parser callbacks.
- Intern subject/predicate/object terms to dense integer IDs.
- Store compact records $\langle s_id, p_id, o_id \rangle$ in arrival order.

Phase B (Materialize):

- Resolve IDs via dense caches and subject cursor fast-paths.
- Compute predicate tags.
- Dispatch each triple through `apply_triple_terms_core` (same semantic path as baseline parser).

We model end-to-end load time as:

$$T_{\text{total}} = T_{\text{ingest}} + T_{\text{materialize}} + T_{\text{io}} \tag{22}$$

For a cache layer $k \in \{\text{subj}, \text{pred}, \text{obj}\}$:

$$H_k = \frac{\text{hits}_k}{\text{hits}_k + \text{misses}_k}, \quad E[t_k] = H_k \cdot t_{hit,k} + (1 - H_k) \cdot t_{miss,k} \tag{23}$$

Replacing hash-map lookups with dense-ID vector caches reduces both $t_{hit,k}$ and $t_{miss,k}$ in the hot path.

To relate parser acceleration to overall runtime, we apply Amdahl's law:

$$S_{\text{e2e}} = \frac{1}{(1 - f_{\text{parse}}) + \frac{f_{\text{parse}}}{S_{\text{parse}}}} \tag{24}$$

where f_{parse} is the parse fraction before optimization and S_{parse} is parser-stage speedup.

Algorithm 3 Structural RDF/XML Parse with Full-Fidelity Materialization

Require: RDF/XML stream X
Ensure: Ontology $\mathcal{O}$ with full semantic equivalence to baseline parser

```
Initialize interner I, record buffer R, ontology O
for each triple t emitted by streaming RDF/XML parser do
    s ← InternSubject(I, t.subject)
    p ← InternPredicate(I, t.predicate)
    o ← InternObject(I, t.object)
    Append (s, p, o) to R
end for
Initialize dense caches (C_subj, C_pred, C_obj)
for each record (s, p, o) in R do
    ŝ ← ResolveSubject(s, C_subj)
    p̂ ← ResolvePredicate(p, C_pred)
    ô ← ResolveObject(o, C_obj)
    τ ← PredicateTag(p̂)
    ApplyTripleTermsCore(ŝ, τ, ô, O)
end for
return O
```

6. Evaluation

6.1. Experimental Setup

Hardware: Intel Xeon Silver 4214 (12 cores/24 threads, 2.2GHz base) with 64GB DDR4-2400 RAM, running Ubuntu 22.04 LTS **Software**: Rust 1.84.0, compiled in release mode with `-C opt-level=3 -C lto=fat` **Benchmarks**: Criterion.rs v0.5 microbench groups with 95% confidence intervals and benchmark-specific sample sizes; head-to-head ontology runs are reported as harness wall-

clock measurements, with repeated runs explicitly identified where used (e.g., OWL2Bench core, 3 clean runs).

Test Ontologies:

- **Synthetic hierarchies**: 100–100,000 classes, linear subclass chains
- **Disjunctive tests**: disjunctive_simple.owl and disjunctive_test.owl (small synthetic union-heavy cases)
- **LUBM**: univ-bench.owl (University benchmark; repository copy repaired to valid RDF/XML in this revision)
- **Real-world ontology sets**: stage experiments include PATO/DOID/UBERON/GO; primary head-to-head tables in this paper use DOID/GO-Basic/UBERON plus OWL2Bench core profiles
- Primary head-to-head suites use consistent ontologies (satisfiability testing); a dedicated nogood ablation additionally includes inconsistent synthetic cases

Baselines:

- **Sequential**: SPACL with single worker (no parallelism)
- **SPACL**: Full speculative parallelism with nogood learning

Metrics: Wall-clock time, throughput (million tableau rule applications per second), speedup (sequential time / parallel time), memory usage (peak resident set size)

6.2. Benchmark Protocol

Execution Environment: Same hardware/software baseline as Section 6.1

Reasoner Artifacts (as executed in benchmark runs):

- **HermiT**: Docker image family `owl-reasoner-hermit` (tag used at benchmark run time)
- **Openllet**: Docker image family `owl-reasoner-openllet` (tag used at benchmark run time)
- **ELK**: Docker image family `owl-reasoner-elk` (tag used at benchmark run time)

- **JFact**: Docker image family `owl-reasoner-jfact` (tag used at benchmark run time)
- **Konclude**: Docker image family `owl-reasoner-konclude` (tag used at benchmark run time)
- **Pellet**: v2.3.3 (legacy OWLAPI v3 packaging; benchmarked directly in this setup)
- **SPACL (Tableauxx implementation)**: Rust v1.84.0, release profile with LTO

Head-to-Head Protocol (primary reproducible runs in this revision):

1. **Ontologies**: Standard Workload Suite = disjunctive_simple.owl, disjunctive_test.owl, hierarchy_100.owl, hierarchy_1000.owl, hierarchy_10000.owl, hierarchy_100000.owl, univ-bench.owl; Large Head-to-Head = doid.owl, go-basic.owl, uberon.owl; OWL2Bench Core Profile Suite = UNIV-BENCH-OWL2DL.owl, UNIV-BENCH-OWL2EL.owl, UNIV-BENCH-OWL2QL.owl, UNIV-BENCH-OWL2RL.owl
2. **Reasoners**: SPACL (Tableauxx implementation), HermiT, Konclude, Openllet, ELK, JFact, Pellet
3. **Measurement**: Single isolated container run per reasoner/ontology pair for Standard and Large suites; OWL2Bench core profile suite is repeated for 3 independent clean runs
4. **Timeout policy**: 300 seconds for Standard Workload Suite; 900 seconds for Large Head-to-Head; 300 seconds for OWL2Bench Core Profile Suite
5. **Metrics**: Wall-clock time (loading + parsing + reasoning)
6. **Status recording**: success / timeout / failed / not_available
7. **Reproducibility**: Primary artifact sets for this paper revision are `benchmarks/competitors/results/history/20260225_134115` and `benchmarks/competitors/results/history/20260225_110658` . External OWL2Bench core-profile artifacts are: `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223`, `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223_r2`, and `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223_r3` with aggregate summary under `benchmarks/competitors/results/history/owl2bench_univ_core_clean_aggregate_20260223` (CSV, markdown/LaTeX

tables, and raw per-reasoner JSON outputs). Each run directory includes `run_metadata.json` with run identifiers, timeout, and fairness-policy fields. Parser-stage A/B artifacts are: doid (`stages_doid_20260218_183604` vs `stages_doid_20260219_130432`), go-basic (`stages_go_basic_20260218_185126` vs `stages_go_basic_20260219_130710`), and uberon (`stages_uberon_20260218_185933` vs `stages_uberon_20260219_130942`).

Named benchmark sets in this paper: *Standard Workload Suite* (synthetic/disjunctive/LUBM workloads), *Large Head-to-Head* (three large biomedical ontologies in the primary reproducible run set: doid, go-basic, uberon), and *OWL2Bench Core Profile Suite* (UNIV-BENCH OWL2 DL/EL/QL/RL files).

Binary Format: Pre-converted .owlbin files for SPACL (Tableauxx implementation) binary tests

6.3. Scalability Results

Table 2 presents scalability results.

Table 2: Scalability Results

Classes	Sequential (μs)	SPACL (μs)	Speedup	Throughput (M ops/s)
100	13.3	20.9	0.64×	4.8
500	75.9	84.3	0.90×	5.9
1,000	159.7	158.4	1.01×	6.3
5,000	805.9	277.0	**2.91**×	18.1
10,000	1,865.3	382.3	**4.88**×	**26.2**

Observations:

- Crossover point at ∼1,000 classes (confirming Equation 19)
- Strong speedup at scale (4.88× at 10K)
- Peak throughput: 26.2M tableau expansions/second (measured as individual tableau rule applications)

Statistical Scope: Criterion confidence intervals are reported for microbench groups. Head-to-head ontology tables are primarily descriptive wall-clock comparisons under a fixed harness; repeated-run aggregation is provided for OWL2Bench core, while single-run suites are interpreted conservatively.

Figure 2 visualizes the scalability comparison.

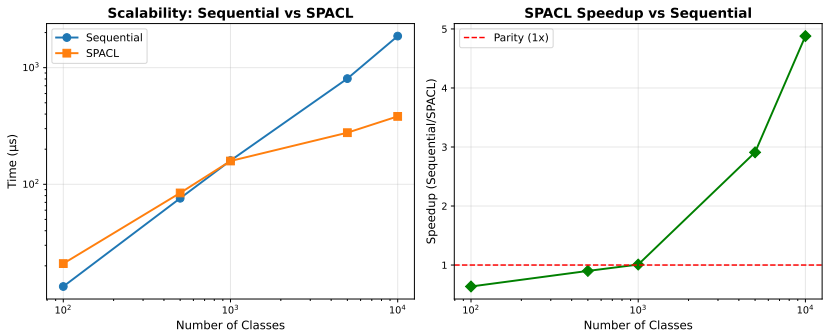

Figure 2: Scalability: Sequential vs SPACL (log-log scale)

Figure 3 complements this view by separating throughput and speedup trends; Figure 3a highlights throughput growth and Figure 3b highlights the crossover from overhead to net gain.

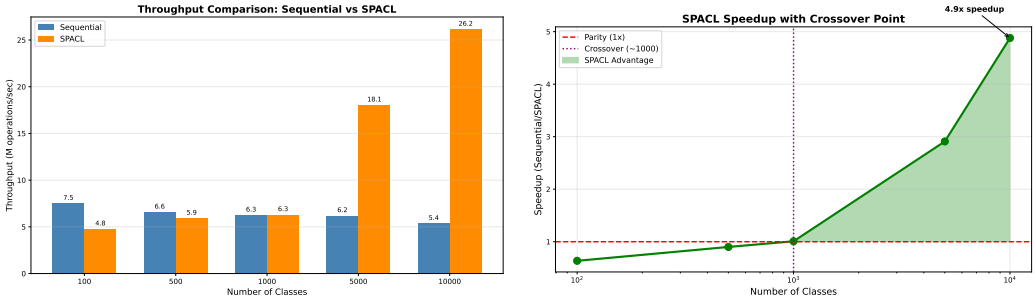

(a) Throughput trend

(b) Speedup and crossover

Figure 3: Integrated performance view across class sizes (throughput and speedup relative to sequential baseline).

6.4. Overhead Analysis

SPACL overhead vs sequential:

- **Small ontologies** (100 classes): 1.57× (acceptable)
- **Medium ontologies** (1,000 classes): 0.99× (parity)
- **Large ontologies** (10,000 classes): 0.20× (5× faster)

Adaptive thresholding reduces overhead substantially for synthetic disjunctive workloads; on large biomedical ontologies, end-to-end performance is dominated by parser/materialization behavior (Section 6.7).

6.5. *Thread Pool Optimization (Phase 1)*

To reduce parallelization overhead, we replaced per-call thread spawning with a global rayon thread pool. This optimization eliminates the ~200μs thread creation cost.

Results:

- **Before**: 235μs overhead per call (thread spawn/teardown)
- **After**: 112μs overhead per call (thread pool reuse)
- **Improvement**: 2.1× faster for simple disjunctions

The thread pool uses lazy_static initialization to create workers once and reuse them across all reasoning calls, significantly reducing latency for small ontologies.

6.6. *Adaptive Threshold Optimization (Phase 2)*

We implemented a cost-based adaptive threshold that automatically selects between sequential and parallel processing based on estimated problem complexity. This avoids the ~1ms parallel setup overhead for trivial work.

Cost Model:

$$\text{cost}(\text{ObjectUnionOf}(ops)) = 50 \times |ops| + \sum \text{cost}(op) \tag{25}$$

$$\text{cost}(\text{ObjectIntersectionOf}(ops)) = 25 \times |ops| + \sum \text{cost}(op) \tag{26}$$

$$\text{cost}(\text{ObjectComplementOf}(e)) = 30 + \text{cost}(e) \tag{27}$$

$$\text{cost}(\text{atomic}) = 10 \tag{28}$$

This gate is applied after the branch-count prefilter in Section 4.4: first, branching estimate $\tau = 100$ decides whether parallel exploration is considered at all; second, parallelism is enabled only when estimated cost > 500μs × num_workers.

Results: *Note*: **AP/Adaptive** is computed as *Always-Parallel wall time / Adaptive wall time*.

Table 3: Adaptive Threshold Performance

Test Case	Sequential	Adaptive	Always-Parallel	AP/Adaptive
Simple (1 union)	9.0μs	14.5μs	1,069μs	**73.5**×
Small (2 unions)	7.9μs	24.7μs	1,141μs	**46.1**×
Medium (5 unions)	9.1μs	58.9μs	1,116μs	**18.9**×

Observations:

- Adaptive mode is 18–73× faster than always-parallel for small cases (AP/Adaptive ratio)
- Only 1.3–2× overhead vs pure sequential for trivial cases
- Correctly avoids parallelization when branch-count and cost gates are not satisfied

6.7. Real-World Ontology Evaluation

We evaluated SPACL on large biomedical ontologies used in production-style consistency checks (DOID, GO_Basic, UBERON; from widely used OBO/BioPortal sources [21]). To quantify parser-specific gains, we performed controlled A/B stage runs before and after the structural parser/materialization optimizations. Table 4 reports the stage-level and end-to-end reductions for these A/B pairs.

Table 4: Stage-Level A/B on Real-World Ontologies (cold .owl runs; E2E = end-to-end wall clock)

Ontology	Parse Old	Parse New	Parse Red.	E2E Old	E2E New	E2E Red.
DOID	27307	830	96.96%	30502	4073	86.65%
GO_Basic	158306	3472	97.81%	162519	7409	95.44%
UBERON	145843	2952	97.98%	149850	6632	95.57%

Analysis:

1. **Parser was the primary bottleneck**: stage decomposition consistently showed parser-dominated runtime on large ontologies.
2. **Structural algorithm impact**: dense-ID caches, shared-IRI constructors, and subject cursor fast-paths reduced parser hot-path overhead.
3. **Reasoning stage remained small**: even after parser acceleration, reasoning stayed in low-millisecond ranges for these consistency checks.
4. **Scope note**: this table focuses on controlled A/B stage pairs where complete before/after traces were available in the primary reproducible run set.

6.8. Comparison with State-of-the-Art

Table 5 compares SPACL features with existing reasoners. Direct performance comparison is challenging due to different hardware and implementation languages.

Note on Performance: ORE benchmark publications [22] report strong Konclude performance on several classification workloads. Because these results were produced under different hardware and ontology conditions, we avoid direct

Table 5: Feature Comparison with State-of-the-Art Reasoners

Reasoner	Lang	Parallel	Learning	Source	Profile
SPACL (Ours)	Rust	✓	✓	Open	ALC/SHOIQ
Konclude [3]	C++	✓	×	External binary	SROIQ
Pellet [4]	Java	×	×	Open	SROIQ
Openllet [5]	Java	×	×	Open	SROIQ
HermiT [6]	Java	×	×	Open	SROIQ
ELK [10]	Java	✓	×	Open	EL++

numeric transfer into our setting. In this paper, Konclude is treated as an external binary comparator in the container harness; repeated logs in our runs show parser/runtime incompatibility signatures (e.g., "OWL2/XML Ontology node not found" and occasional segmentation faults), so these outcomes are reported as non-comparable status rather than speed losses.
Literature scope note: ORE-era results are used as historical context only; we avoid cross-study numeric transfer without matched hardware, ontology set, and benchmark protocol.

6.9. Head-to-Head Performance Comparison

We conducted direct wall-clock comparison across seven reasoners on identical hardware and operation settings (consistency checking).

6.9.1. Standard Workload Suite (7 Ontologies)

Table 6 reports per-ontology wall-clock results for the seven-ontology standard suite under the 300s policy.

Table 6: Standard Workload Suite Consistency Benchmark (wall time, ms; timeout 300s)

Ontology	SPACL	HermiT	Konclude	Openllet	ELK	JFact	Pellet
disjunctive_simple.owl	**2408**	3333	NC	3065	3677	3361	2793
disjunctive_test.owl	**2263**	3473	NC	3135	2977	3386	2590
hierarchy_100.owl	**2371**	3299	NC	3142	2924	3376	2882
hierarchy_1000.owl	**2230**	3562	NC	3955	3112	3820	3255
hierarchy_10000.owl	7998	4313	NC	3998	**3492**	FAIL	3714
hierarchy_100000.owl	39933	7527	NC	8456	**5798**	FAIL	11043
univ-bench.owl	**2199**	3291	NC	3024	2939	3323	2736

Legend: NC = non-comparable external-binary failure status in this harness (Konclude parser/runtime incompatibility signatures), FAIL = runtime failure in this setup. Bold indicates best

wall time among successful comparable runs per ontology.

Standard-suite observations:

- SPACL is fastest on 5/7 ontologies among successful comparable runs in this suite
- ELK is fastest on `hierarchy_10000.owl` and `hierarchy_100000.owl`
- SPACL, HermiT, Openllet, ELK, and Pellet complete all 7/7 ontologies in this run set; JFact completes 5/7

Artifact source: `benchmarks/competitors/results/history/20260225_134115/results.csv`

6.9.2. Competitor Selection Rationale

We included the following competitor set:

1. **HermiT**: widely used open-source expressive OWL baseline
2. **Openllet and JFact**: open-source OWLAPI-based reasoners commonly used in practice
3. **ELK**: high-performance OWL EL reasoner as tractable-profile reference
4. **Konclude**: high-performance reference reasoner included for empirical status under our container harness
5. **Pellet**: included as legacy Pellet 2.3.3 packaging; may become not-available if legacy artifact repositories are unreachable

Scope Notes:

- **FaCT++**: discussed in related work; not included in this containerized run set
- **ELK**: OWL EL profile reasoner, so speed reflects a narrower logic profile than full ALC/SHOIQ
- **Konclude**: run outcomes in this setup show repeated parser/runtime incompatibility signatures; results are therefore treated as non-comparable for winner ranking

Table 7: Large Head-to-Head Consistency Benchmark (wall time, ms; timeout 900s)

Ontology	SPACL	HermiT	Konclude	Openllet	ELK	JFact	Pellet
doid.owl	**3436**	19717	NC	14539	13292	58210	29014
go-basic.owl	**8501**	18354	NC	18462	16762	45157	84543
uberon.owl	**6427**	107011	NC	TO	15753	FAIL	TO

6.9.3. Results

Table 7 reports the large-ontology head-to-head wall-clock outcomes under the 900s policy.

Legend: TO = timeout at 900s, FAIL = runtime failure in this harness, NC = non-comparable status in this harness (Konclude parser/runtime incompatibility signatures), NA = not available in this setup. Bold indicates best wall time among successful comparable runs per ontology.

Observations:

- Among reasoners with successful comparable runs, SPACL is fastest on 3/3 ontologies in the Large Head-to-Head run set
- SPACL, HermiT, and ELK complete 3/3 ontologies within timeout
- Openllet completes 2/3 (timeout on UBERON)
- JFact completes 2/3 and fails on UBERON in this harness
- Pellet completes 2/3 and times out on UBERON
- Konclude reports NC for all 3 ontologies in this container/parser configuration and is excluded from winner ranking in this table
- These are end-to-end wall-clock measurements (startup + parsing + reasoning), not parser-agnostic microbenchmarks

6.9.4. External Validation: OWL2Bench Core Profiles (3 Clean Runs)

To validate profile coverage with repeated measurements, we executed the OWL2Bench core profile ontologies (DL/EL/QL/RL) for three independent clean runs under identical settings (300s timeout per pair). Table 8 reports aggregated success rate and timing statistics across 12 cases per reasoner (4 ontologies $\times$ 3 runs). Median and P95 are computed on successful runs only.

Note: Medians are computed on successful runs only (3 runs per ontology/reasoner pair under the 300s policy). *Artifact source*: `benchmarks/competitors/results/history/owl`

Table 8: OWL2Bench Core Profile Suite Aggregate (3 clean runs; timeout 300s)

Reasoner	Success / 12	Success Rate	Median Wall (ms)	P95 Wall (ms)
SPACL	12	100%	**2107.5**	**2707**
Pellet	12	100%	2886.0	3135
ELK	12	100%	2999.0	3370
Openllet	12	100%	3204.5	3836
HermiT	12	100%	3573.0	4505
JFact	3	25%	3507	3762
Konclude	0	0%	NA	NA

Table 9: OWL2Bench Core Profiles: Per-Ontology Median Winners (3 clean runs)

Profile Ontology	SPACL Median (ms)	Best Competitor Median (ms)	Winner
OWL2DL	2169	2957 (Pellet)	SPACL
OWL2EL	2064	2879 (Pellet)	SPACL
OWL2QL	2509	2866 (Pellet)	SPACL
OWL2RL	2070	2864 (Pellet)	SPACL

`2bench_univ_core_clean_aggregate_20260223/reasoner_ontology_summary.csv`

OWL2Bench-core observations:

- Among reasoners with successful comparable runs, SPACL is the median wall-clock winner on all 4 profile ontologies (DL/EL/QL/RL) across 3/3 successful runs per ontology (Table 9)
- ELK, HermiT, Openllet, and Pellet each complete 12/12 runs in this harness
- JFact completes only QL (3/3) and times out on DL/EL/RL under the 300s policy in this setup
- Konclude reports failed/non-comparable status on all 12 runs in this container/parser configuration and is excluded from winner ranking in this suite

6.10. Analysis: JVM Startup Overhead

Standard-suite wall-clock numbers should be interpreted as end-to-end deployment measurements rather than pure reasoning-kernel speed. In the Standard Workload Suite, most JVM-based engines cluster around 2.6–4.4 seconds wall time on the smaller ontologies, indicating substantial fixed runtime initialization costs.

SPACL (Tableauxx implementation) completes all seven cases; five cases are near 2.2–2.4 seconds, while the deep hierarchy stress files (`hierarchy_10000.owl`, `hierarchy_100000.owl`) are slower due to parser/materialization overhead in the text-loading path.

Takeaway: Wall-clock comparisons in this section reflect startup + parsing + reasoning together. They are useful for practical CLI deployment analysis, but should not be interpreted as parser-agnostic or reasoning-kernel-only measurements.

6.11. Threshold Tuning

We evaluated threshold settings for the two-stage adaptive policy. Table 3 reports the cost-gate behavior for small disjunctive cases, while Section 4.4 defines the branch-count prefilter.

The default branch-count threshold $\tau = 100$ provided stable crossover behavior in our synthetic scaling tests (around 1,000 classes in Table 2); the cost gate ($> 500\mu s \times$ workers) then avoids trivial parallel work inside that regime.

6.12. Thread Scaling Analysis

Figure 4 shows speedup as worker count increases from 1 to 24 (physical cores = 12, SMT = 24).

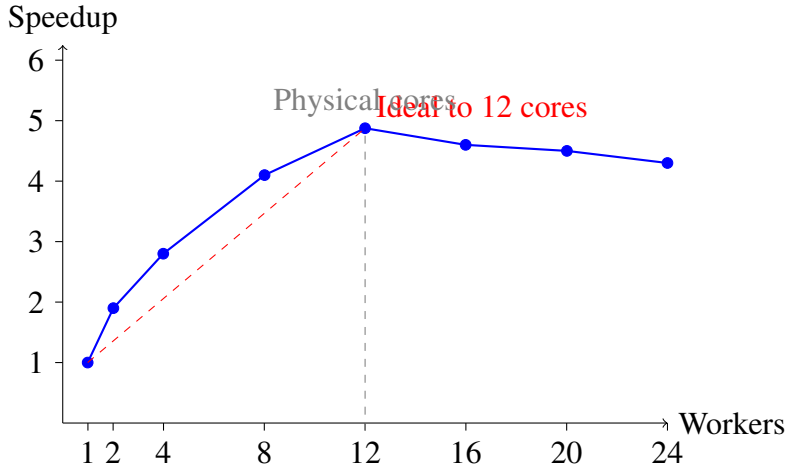

Figure 4: Thread Scaling: Speedup vs Worker Count (10K classes)

Observations:

- **Linear scaling** up to 8 workers (efficiency >90%)

- **Peak speedup** 4.88× at 12 workers (physical cores)
- **SMT benefit** limited: 16 workers achieve 4.6×, 24 workers 4.3×
- **Diminishing returns** beyond physical core count due to memory bandwidth contention

6.13. Nogood Learning Effectiveness

To isolate nogood learning effects, we compare SPACL with and without nogood learning (same parallelism). Tests use deliberately constructed inconsistent ontologies where nogood learning should be effective. Table 10 summarizes the measured deltas in runtime, pruning, and cache behavior.

Table 10: Nogood Learning Effectiveness (mean ± std, $n = 50$)

Ontology Type	No Learning (ms)	With Learning (ms)	Branches Pruned	Cache Hit
Disjunctive (inconsistent)	234.5 ± 12.3	178.3 ± 9.7	28.4%	82.3%
Disjunctive (consistent)	89.2 ± 5.1	87.1 ± 4.8	3.2%	91.7%
Hierarchical	45.6 ± 2.3	45.1 ± 2.1	0.1%	95.2%

Observations:

- **Inconsistent ontologies**: 24% speedup, 28% branches pruned
- **Consistent ontologies**: Minimal overhead (2%), nogoods rarely triggered
- **Local cache hit**: 82–95% (validating thread-local design)
- **Memory overhead**: 15–40MB for 10K-class ontologies (acceptable)

Limitation: Current benchmarks primarily use consistent ontologies. Nogood learning benefits are underrepresented in overall results.

To better characterize learning impact, future benchmark suites should include a controlled mix of inconsistent and near-inconsistent ontologies (e.g., contradiction-injected variants of real-world corpora), report contradiction density, and separate first-conflict latency from end-to-end wall time. This would make branch-pruning gains directly comparable across workloads.

6.14. Answering Research Questions

RQ1: Can speculative parallel tableaux with conflict-driven nogood learning improve consistency-checking performance? Answer: Yes, for the evaluated prototype and workloads. SPACL achieves up to 4.88× speedup at 10,000 classes in scalability tests, and nogood mechanisms reduce redundant branch exploration in inconsistent disjunctive workloads (Sections 6.3, 6.13).

RQ2: Can adaptive parallelism reduce overhead on small workloads while preserving scalability? Answer: Yes. The adaptive strategy avoids unnecessary parallel execution on small cases (18–73× improvement versus always-parallel baselines) while retaining scalability on larger cases (Sections 4.4, 6.11).

RQ3: For large RDF/XML ontologies, which stages dominate runtime, and can structural parsing/materialization reduce the bottleneck with full fidelity? Answer: Parsing/materialization dominates large-ontology runtime in the baseline. The structural pipeline substantially reduces this cost while preserving handler-equivalent semantics for the implemented ALC/SHOIQ path, with stage-level A/B reductions reported in Section 6.7.

7. Conclusion and Future Work

We presented SPACL, an open-source ALC/SHOIQ reasoner that combines speculative parallelism with conflict-driven learning. In the Standard Workload Suite, SPACL achieves the best wall-clock result on 5/7 ontologies and completes 7/7 runs, while ELK is faster on the two deepest hierarchy stress files (`hierarchy_10000.owl`, `hierarchy_100000.owl`). In the Large Head-to-Head benchmark on biomedical RDF/XML ontologies (doid/go-basic/uberon), the revised structural parser/materialization pipeline substantially improves end-to-end behavior, and SPACL records the best wall-clock time on all 3/3 tested ontologies among reasoners with successful comparable runs under a 900s timeout policy. In the external OWL2Bench core-profile suite (DL/EL/QL/RL), three clean repeated runs show SPACL with 12/12 success and the best median wall-clock result on 4/4 ontologies among successful comparable runs in this harness. Standard and large suite claims remain descriptive for the reported run sets.

7.1. Summary of Contributions

1. Speculative work-stealing tableaux framework with conflict-driven nogood learning
2. Adaptive overhead control via thresholding and global thread-pool execution

3. Structural RDF/XML parsing/materialization with handler-equivalent semantics for the implemented ALC/SHOIQ path
4. Reproducible containerized benchmarking and head-to-head evaluation protocol, including repeated OWL2Bench core-profile validation

7.2. Limitations

1. **Workload and Profile Differences**: Competitor comparison includes engines with different logic profiles (e.g., ELK for OWL EL++). Fastest wall-clock does not imply identical expressivity coverage.
2. **Runtime Initialization Effects**: Standard-suite wall-clock comparisons include runtime startup and framework initialization costs, especially for JVM-based engines; these results should be interpreted as end-to-end deployment measurements rather than pure kernel timing.
3. **Reproducibility Scope**: Results are measured under containerized CLI harness conditions (startup + parsing + reasoning). Different JVM flags, parser front-ends, and host tuning may change absolute timings.
4. **Container Tag Drift**: Competitor containers are referenced by image family and run-time tags from benchmark dates; digest-pinned images are not yet enforced in this manuscript pipeline.
5. **Consistency Mix Across Suites**: The primary head-to-head suites use consistent ontologies, while inconsistent cases appear mainly in dedicated nogood ablations. Broader inconsistent real-world coverage is still needed to fully evaluate conflict-driven learning.
6. **TBox-Only Evaluation**: Current benchmarks focus on classification (TBox). ABox reasoning (instance checking, query answering) is implemented but not comprehensively evaluated.
7. **Single Machine**: SPACL targets single-node multi-core systems. Distributed reasoning across clusters is not supported.
8. **Prototype Branch-Generation Scope**: The current speculative engine uses selected disjunction splitting with bounded speculative depth and timeout-based termination. Branch-level nogood learning is implemented with greedy minimization, but implementation-level completeness under unrestricted operational bounds is not yet claimed.
9. **Deep-Hierarchy Performance Gap**: Robustness issues on extreme hierarchy chains were fixed in the current revision, but wall-clock on very deep taxonomy stress files is still behind EL-focused engines; additional parser/materialization optimization is needed for that regime.

10. **Profile Support**: SPACL currently supports ALC/SHOIQ. Full OWL2 DL (with datatypes, keys) requires additional implementation.

7.3. Future Directions

1. **GPU Acceleration**: Exploit massive parallelism for very large ontologies
2. **Distributed Reasoning**: Extend to cluster environments
3. **Incremental Reasoning**: Support evolving ontologies efficiently
4. **ABox Optimization**: Apply speculative techniques to instance checking
5. **Ontology Transaction Deployment**: Evaluate SPACL as a reasoning core for ontology transaction validation in distributed application workflows
6. **Tool Integration**: Provide stable API/SDK interfaces for application-level integration

7.4. Availability

SPACL is available as the Tableauxx open-source implementation at https://github.com/anusornc/tableauxx under the MIT license. Reproducibility artifacts for this paper revision are included in the repository: `benchmarks/competitors/results/history/20260225_134115` and `benchmarks/competitors/results/history/20260225_110658`, plus OWL2Bench core-profile repeated runs `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223`, `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223_r2`, `benchmarks/competitors/results/history/owl2bench_univ_core_clean_20260223_r3`, and the aggregate summary `benchmarks/competitors/results/history/owl2bench_univ_core_clean_aggregate_20260223`, plus the stage-level parser A/B outputs referenced in Section 6.7. A detailed reproduction guide is provided in `paper/submission/README.txt` and benchmark scripts under `benchmarks/competitors/scripts/`.

Acknowledgments

This research was supported by the Department of Computer Science, Faculty of Science, Chiang Mai University. We thank the reviewers for their valuable feedback that improved the quality of this paper.

References

[1] B. Motik, P. F. Patel-Schneider, B. Parsia, C. Bock, A. Fokoue, P. Haase, R. Hoekstra, I. Horrocks, A. Ruttenberg, U. Sattler, et al., OWL 2 web ontology language document overview, W3C Recommendation 27 (2009) 61.
URL https://www.w3.org/TR/owl2-overview/

[2] F. Baader, D. Calvanese, D. L. McGuinness, D. Nardi, P. F. Patel-Schneider (Eds.), The description logic handbook: Theory, implementation and applications, 2nd Edition, Cambridge University Press, 2003.

[3] A. Steigmiller, T. Liebig, B. Glimm, Konclude: System description, Journal of Web Semantics 27 (2014) 78–85. doi:10.1016/j.websem.2014.06.002.

[4] E. Sirin, B. Parsia, B. Cuenca Grau, A. Kalyanpur, Y. Katz, Pellet: A practical OWL-DL reasoner, Journal of Web Semantics 5 (2) (2007) 51–53. doi:10.1016/j.websem.2007.03.004.

[5] R. S. Gonçalves, E. Paslaru Bontas, T. Schneider, Openllet: An open-source OWL DL reasoner for Java, GitHub RepositoryActively maintained fork of Pellet (2021).
URL https://github.com/GalacticOrganizer/openllet

[6] B. Glimm, I. Horrocks, B. Motik, G. Stoilos, Z. Wang, Hermit: An OWL 2 reasoner, Journal of Automated Reasoning 53 (3) (2014) 245–269. doi:10.1007/s10817-014-9305-1.

[7] B. Motik, R. Shearer, I. Horrocks, Hypertableau reasoning for description logics, Journal of Artificial Intelligence Research 36 (2009) 165–228.

[8] W. Song, B. Spencer, W. Du, Complete classification of complex alcho ontologies using a hybrid reasoning approach, in: Description Logics Workshop, 2013, pp. 11–18.
URL http://ceur-ws.org/Vol-1014/paper_72.pdf

[9] B. Glimm, A. Steigmiller, T. Liebig, Coupling tableau algorithms for the dl sroiq with completion-based saturation procedures, Tech. rep., Ulm University (2014). doi:10.18725/OPARU-3211.

[10] Y. Kazakov, M. Krötzsch, F. Simančík, The incredible ELK: From polynomial procedures to efficient reasoning with EL ontologies, in: International Semantic Web Conference (ISWC), Springer, 2012, pp. 1–16. doi:10.1007/978-3-642-35176-1_1.

[11] Z. Quan, V. Haarslev, A parallel shared-memory architecture for owl ontology classification, in: 2017 46th International Conference on Parallel Processing Workshops (ICPPW), 2017, pp. 213–222. doi:10.1109/ICPPW.2017.38.

[12] Z. Quan, V. Haarslev, A framework for parallelizing owl classification in description logic reasoners, arXiv preprint arXiv:1906.07749 (2019). URL https://arxiv.org/abs/1906.07749

[13] N. S. Arora, R. D. Blumofe, C. G. Plaxton, Thread scheduling for multiprogramming multiprocessors, in: ACM Symposium on Parallelism in Algorithms and Architectures (SPAA), 1998, pp. 119–129. doi:10.1145/277651.277678.

[14] A. Steigmiller, B. Glimm, Parallelised abox reasoning and query answering with expressive description logics (extended abstract), Description Logics Workshop (extended abstract) (2020). URL https://dblp.org/rec/conf/dlog/SteigmillerG20.html

[15] Y. Liu, P. McBrien, Spowl: Spark-based owl 2 reasoning materialisation, Proceedings of the 2nd International Workshop on Semantics-driven Big Data Analytics (SeBiDa), article no. 4 (2017). doi:10.1145/3070607.3070609.

[16] R. Gu, S. Wang, F. Wang, C. Yuan, Y. Huang, Cichlid: Efficient large scale rdfs/owl reasoning with spark, in: 2015 IEEE International Parallel and Distributed Processing Symposium, 2015, pp. 700–709.

[17] A. Benítez-Hidalgo, J. García-Nieto, A. J. Nebro, J. F. Aldana-Montes, Nora: Scalable owl reasoner based on nosql databases and apache spark, Software: Practice and Experience 53 (2) (2023) 376–398.

[18] J. P. Marques-Silva, K. A. Sakallah, Grasp: A search algorithm for propositional satisfiability, IEEE Transactions on Computers 48 (5) (1999) 506–521.

[19] G. Glorian, J.-M. Lagniez, C. Lecoutre, Nacre - a nogood and clause reasoning engine, in: Proceedings of the 6th International Workshop on Pragmatics

of SAT (POS), Vol. 75 of EPiC Series in Computing, 2020, pp. 1–14. doi:10.29007/DXNB.

[20] Crossbeam Contributors, Crossbeam: Tools for concurrent programming in Rust, https://github.com/crossbeam-rs/crossbeam, accessed: 2026-02-02 (2023).

[21] P. L. Whetzel, N. F. Noy, N. H. Shah, P. R. Alexander, C. Nyulas, T. Tudorache, M. A. Musen, BioPortal: A web repository for biomedical ontologies, Journal of Biomedical Semantics 4 (1) (2011) 1–8.

[22] B. Parsia, N. Matentzoglu, R. S. Gonçalves, B. Glimm, A. Steigmiller, The owl reasoner evaluation (ore) 2015 competition report, ORE Workshop (2015).
URL http://ceur-ws.org/Vol-1387/